



MANIPULATING TIME SERIES DATA IN PYTHON

How to use Dates & Times with pandas

Date & Time Series Functionality

- At the root: data types for date & time information
 - Objects for points in time and periods
 - Attributes & methods reflect time-related details
- Sequences of dates & periods:
 - Series or DataFrame columns
 - Index: convert object into Time Series
- Many Series/DataFrame methods rely on time information in the index to provide time-series functionality



Basic Building Block: `pd.Timestamp`

```
In [1]: import pandas as pd # assumed imported going forward
```

```
In [2]: from datetime import datetime # To manually create dates
```

```
In [3]: time_stamp = pd.Timestamp(datetime(2017, 1, 1))
```

```
In [4]: In [13]: pd.Timestamp('2017-01-01') == time_stamp
```

```
Out[4]: True # Understands dates as strings
```

```
In [5]: time_stamp # type: pandas.tseries.Timestamp
```

```
Out[5]: Timestamp('2017-01-01 00:00:00')
```

```
In [6]: time_stamp.year
```

```
Out[6]: 2017
```

```
In [7]: time_stamp.weekday_name
```

```
Out[7]: 'Sunday'
```

Timestamp object has many attributes to store time-specific information



More building blocks: `pd.Period` & `freq`

```
In [8]: period = pd.Period('2017-01')
```

```
In [9]: period # default: month-end  
Out[9]: Period('2017-01', 'M')
```

Period object has `freq` attribute to store frequency info

```
In [10]: period.asfreq('D') # convert to daily  
Out[10]: Period('2017-01-31', 'D')
```

```
In [11]: period.to_timestamp().to_period('M')  
Out[11]: Period('2017-01', 'M')
```

Convert `pd.Period()` to `pd.Timestamp()` and back

```
In [12]: period + 2  
Out[12]: Period('2017-03', 'M')
```

Frequency info enables basic date arithmetic

```
In [13]: pd.Timestamp('2017-01-31', 'M') + 1  
Out[13]: Timestamp('2017-02-28 00:00:00', freq='M')
```

Sequences of Dates & Times

```
In [14]: index = pd.date_range(start='2017-1-1', periods=12, freq='M')
```

pd.date_range: start, end, periods, freq

```
In [15]: index
DatetimeIndex(['2017-01-31', '2017-02-28', '2017-03-31', ...,
              '2017-09-30', '2017-10-31', '2017-11-30', '2017-12-31'],
              dtype='datetime64[ns]', freq='M')
```

```
In [16]: index[0]
Timestamp('2017-01-31 00:00:00', freq='M')
```

pd.DateTimeIndex:
sequence of Timestamp
objects with frequency info

```
In [17]: index.to_period()
PeriodIndex(['2017-01', '2017-02', '2017-03', '2017-04', ...,
            '2017-11', '2017-12'], dtype='period[M]', freq='M')
```



Create a Time Series: `pd.DatetimeIndex`

```
In [14]: pd.DataFrame({'data': index}).info()
```

```
RangeIndex: 12 entries, 0 to 11
```

```
Data columns (total 1 columns):
```

```
data      12 non-null datetime64[ns]
```

```
dtypes: datetime64[ns](1)
```

```
In [15]: data = np.random.random((size=12,2))
```

`np.random.random:`
Random numbers [0,1]
12 rows, 2 columns

```
In [16]: pd.DataFrame(data=data, index=index).info()
```

```
DatetimeIndex: 12 entries, 2017-01-31 to 2017-12-31
```

```
Freq: M
```

```
Data columns (total 2 columns):
```

```
0      12 non-null float64
```

```
1      12 non-null float64
```

```
dtypes: float64(2)
```



Frequency Aliases & Time Info

There are many frequency aliases besides 'M' and 'D':

Period	Alias
Hour	H
Day	D
Week	W
Month	M
Quarter	Q
Year	A

These may be further differentiated by beginning/end of period, or business-specific definition

You can also access these `pd.Timestamp()` attributes:

attribute
<code>.second</code> , <code>.minute</code> , <code>.hour</code> , <code>.day</code> , <code>.month</code> , <code>.quarter</code> , <code>.year</code>
<code>.weekday</code>
<code>dayofweek</code>
<code>.weekofyear</code>
<code>.dayofyear</code>



MANIPULATING TIME SERIES DATA IN PYTHON

Let's practice!



MANIPULATING TIME SERIES DATA IN PYTHON

Indexing & Resampling Time Series

Time Series Transformation

Basic Time Series transformations include:

- Parsing string dates and convert to `datetime64`
- Selecting & slicing for specific subperiods
- Setting & changing `DateTimeIndex` frequency
 - Upsampling vs Downsampling



Getting GOOG stock prices

```
In [1]: google = pd.read_csv('google.csv') # import pandas as pd
```

```
In [2]: google.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 504 entries, 0 to 503
Data columns (total 2 columns):
date      504 non-null object
price     504 non-null float64
dtypes: float64(1), object(1)
```

```
In [3]: google.head()
```

	date	price
0	2015-01-02	524.81
1	2015-01-05	513.87
2	2015-01-06	501.96
3	2015-01-07	501.10
4	2015-01-08	502.68

Converting `string` dates to `datetime64`

```
In [4]: google.date = pd.to_datetime(google.date)
```

```
In [5]: google.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 504 entries, 0 to 503
Data columns (total 2 columns):
date      504 non-null datetime64[ns]
price     504 non-null float64
dtypes: datetime64[ns](1), float64(1)
```

```
In [6]: google.set_index('date', inplace=True)
```

```
In [7]: google.info()
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 504 entries, 2015-01-02 to 2016-12-30
Data columns (total 1 columns):
price     504 non-null float64
dtypes: float64(1)
```

`pd.to_datetime()`:

- Parse date string
- Convert to `datetime64`

`.set_index()`:


- Date into index
- `inplace`:**
- don't create copy



Plotting the Google Stock Time Series

```
In [8]: google.price.plot(title='Google Stock Price')
```

```
In [9]: plt.tight_layout(); plt.show()
```





Partial String Indexing

```
In [10]: google['2015'].info() # Pass string for part of date
```

```
DatetimeIndex: 252 entries, 2015-01-02 to 2015-12-31
```

```
Data columns (total 1 columns):
```

```
price      252 non-null float64
```

```
dtypes: float64(1)
```

**Selecting/indexing using
strings that parse to dates**

```
In [11]: google['2015-3': '2016-2'].info() # Slice includes last month
```

```
DatetimeIndex: 252 entries, 2015-03-02 to 2016-02-29
```

```
Data columns (total 1 columns):
```

```
price      252 non-null float64
```

```
dtypes: float64(1)
```

```
memory usage: 3.9 KB
```

```
In [12]: google.loc['2016-6-1', 'price'] # Use full date with .loc[]
```

```
Out[12]: 734.15
```



.asfreq(): Set Frequency

```
In [13]: google.asfreq('D').info() # set calendar day frequency
```

```
DatetimeIndex: 729 entries, 2015-01-02 to 2016-12-30
```

```
Freq: D
```

```
Data columns (total 1 columns):
```

```
price      504 non-null float64
```

```
dtypes: float64(1)
```

```
In [14]: google.asfreq('D').head()
```

```
Out[14]:
```

	price
date	
2015-01-02	524.81
2015-01-03	NaN
2015-01-04	NaN
2015-01-05	513.87
2015-01-06	501.96

.asfreq('D'):
Convert DateTimeIndex to
calendar day frequency

Upsampling:
Higher frequency implies new
dates => missing data



• `asfreq()`: Reset Frequency

```
In [18]: google = google.asfreq('B') # Change to calendar day frequency
```

```
In [19]: google.info()
```

```
DatetimeIndex: 521 entries, 2015-01-02 to 2016-12-30
```

```
Freq: B
```

```
Data columns (total 1 columns):
```

```
price    504 non-null float64
```

```
dtypes: float64(1)
```

`.asfreq('B')`:
**Convert `DatetimeIndex` to
business day frequency**

```
In [20]: google[google.price.isnull()] # Select missing 'price' values
```

```
Out[20]:
```

	price
date	
2015-01-19	NaN
2015-02-16	NaN
..	
2016-11-24	NaN
2016-12-26	NaN

**Business days that were
not trading days**



MANIPULATING TIME SERIES DATA IN PYTHON

Let's practice!



MANIPULATING TIME SERIES DATA IN PYTHON

Lags, changes, and returns for Stock Price Series

Basic Time Series Calculations

- Typical Time Series manipulations include:
 - Shift or lag values back or forward back in time
 - Get the difference in value for a given time period
 - Compute the percent change over any number of periods
- pandas built-in methods rely on `pd.DatetimeIndex`



Getting GOOG stock prices

```
In [1]: google = pd.read_csv('google.csv',  
                             parse_dates=['date'],  
                             index_col='date')
```

Let `pd.read_csv()` do the parsing for you!

```
In [2]: google.info()  
<class 'pandas.core.frame.DataFrame'>  
DatetimeIndex: 504 entries, 2015-01-02 to 2016-12-30  
Data columns (total 1 columns):  
price      504 non-null float64  
dtypes: float64(1)
```

```
In [3]: google.head()  
           price
```

```
date  
2015-01-02    524.81  
2015-01-05    513.87  
2015-01-06    501.96  
2015-01-07    501.10  
2015-01-08    502.68
```

.shift(): Moving data between past & future

```
In [4]: google['shifted'] = google.price.shift() # default: periods=1
```

```
In [5]: google.head(3)
```

```
Out[5]:
```

	price	shifted
date		
2015-01-02	542.81	NaN
2015-01-05	513.87	542.81
2015-01-06	501.96	513.87

.shift():
defaults to **periods=1**
1 period into future

```
In [6]: google['lagged'] = google.price.shift(periods=-1)
```

```
In [7]: google[['price', 'lagged', 'shifted']].tail(3)
```

```
Out[7]:
```

	price	lagged	shifted
date			
2016-12-28	785.05	782.79	791.55
2016-12-29	782.79	771.82	785.05
2016-12-30	771.82	NaN	782.79

.shift(periods=-1):
lagged data:
1 period back in time

Calculate one-period percent change

```
In [10]: google['change'] = google.price.div(google.shifted) #  $x_t / x_{t-1}$ 
```

```
In [11]: google[['price', 'shifted', 'change']].head(3)
```

```
Out[11]:
```

	price	shifted	change
Date			
2017-01-03	786.14	NaN	NaN
2017-01-04	786.90	786.14	1.000967
2017-01-05	794.02	786.90	1.009048

```
In [12]: google['return'] = google.change.sub(1).mul(100)
```

```
In [13]: google[['price', 'shifted', 'change', 'return']].head(3)
```

```
Out[13]:
```

	price	shifted	change	return
date				
2015-01-02	524.81	NaN	NaN	NaN
2015-01-05	513.87	524.81	0.98	-2.08
2015-01-06	501.96	513.87	0.98	-2.32

`.diff()` & `.pct_change()`: built-in time-series change

```
In [14]: google['diff'] = google.price.diff() #  $x_t - x_{t-1}$ 
```

```
In [15]: google[['price', 'diff']].head(3)
```

```
Out[15]:
```

	price	diff
date		
2015-01-02	524.81	NaN
2015-01-05	513.87	-10.94
2015-01-06	501.96	-11.91

Difference in value for two adjacent periods

```
In [16]: google['pct_change'] = google.price.pct_change().mul(100)
```

```
In [17]: google[['price', 'return', 'pct_change']].head(3)
```

```
Out[17]:
```

	price	return	pct_change
date			
2015-01-02	524.81	NaN	NaN
2015-01-05	513.87	-2.08	-2.08
2015-01-06	501.96	-2.32	-2.32

Percent change for two adjacent periods



Looking ahead: Get Multi-period returns

```
In [25]: google['return_3d'] = google.price.pct_change(3).mul(100)
```

```
In [34]: google[['price', 'return_3d']].head()
```

```
Out[34]:
```

	price	return_3d
date		
2015-01-02	524.81	NaN
2015-01-05	513.87	NaN
2015-01-06	501.96	NaN
2015-01-07	501.10	-4.517825
2015-01-08	502.68	-2.177594

Percent change for two periods, 3 trading days apart



MANIPULATING TIME SERIES DATA IN PYTHON

Let's practice!